

CONCURRENCY I

Instructors: Crista Lopes
Copyright © Instructors.



Basics

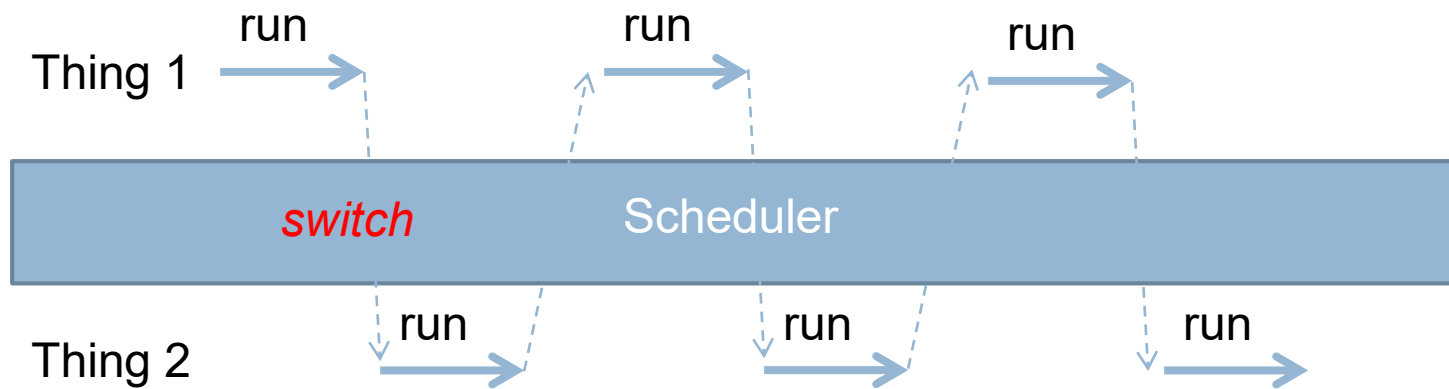


Concurrent Programming

- More than one thing at a time
- Examples:
 - ▣ Network server handling hundreds of clients
 - ▣ Data processing spread across multiple CPUs
 - ▣ App receiving input from several peripherals

Concurrency

□ 1 CPU




Concurrency

- Many CPUs
- When there are more things than CPUs, CPUs are shared like previous slide



Processing vs. IO

- Tasks alternate between processing and IO



- IO must wait for data
- Runtime system will resume task when data becomes available

CPU-bound tasks

- When it spends most of the time processing
 - CPU is busy



- Examples:
 - Math
 - Image processing

IO-bound tasks

- When it spends most of the time doing IO
 - CPU is idle

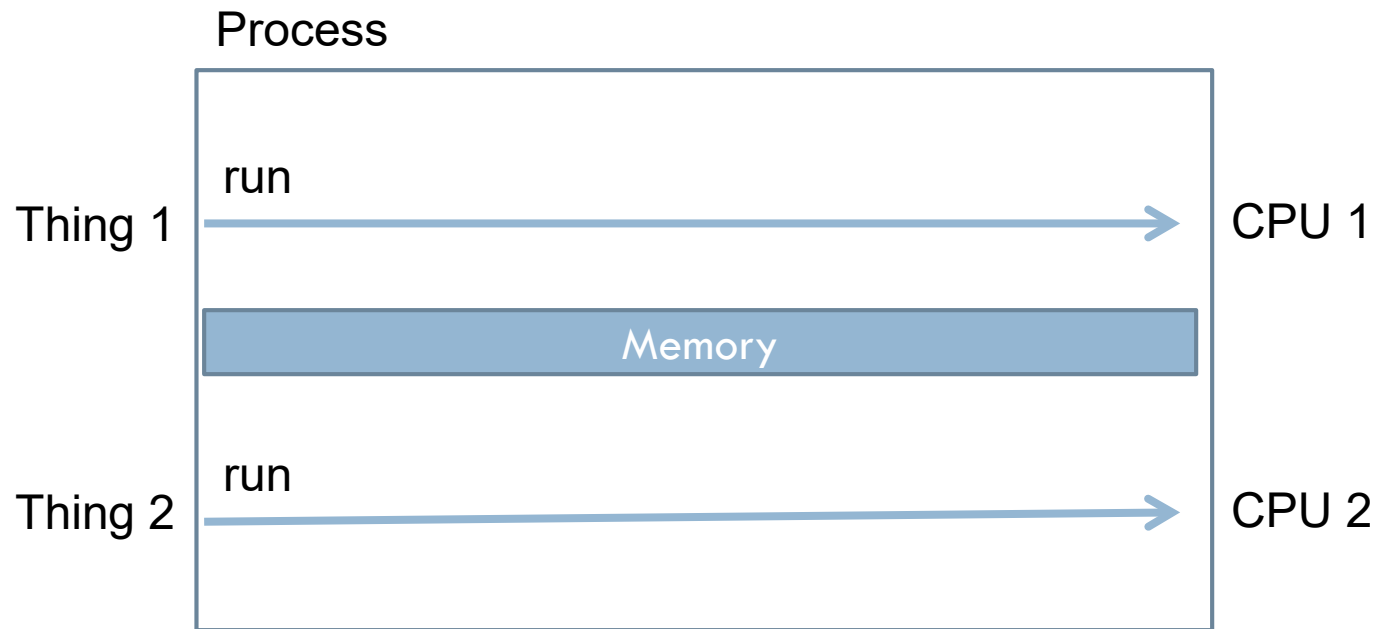


- Examples:
 - User input
 - Networking
 - Files

Shared memory

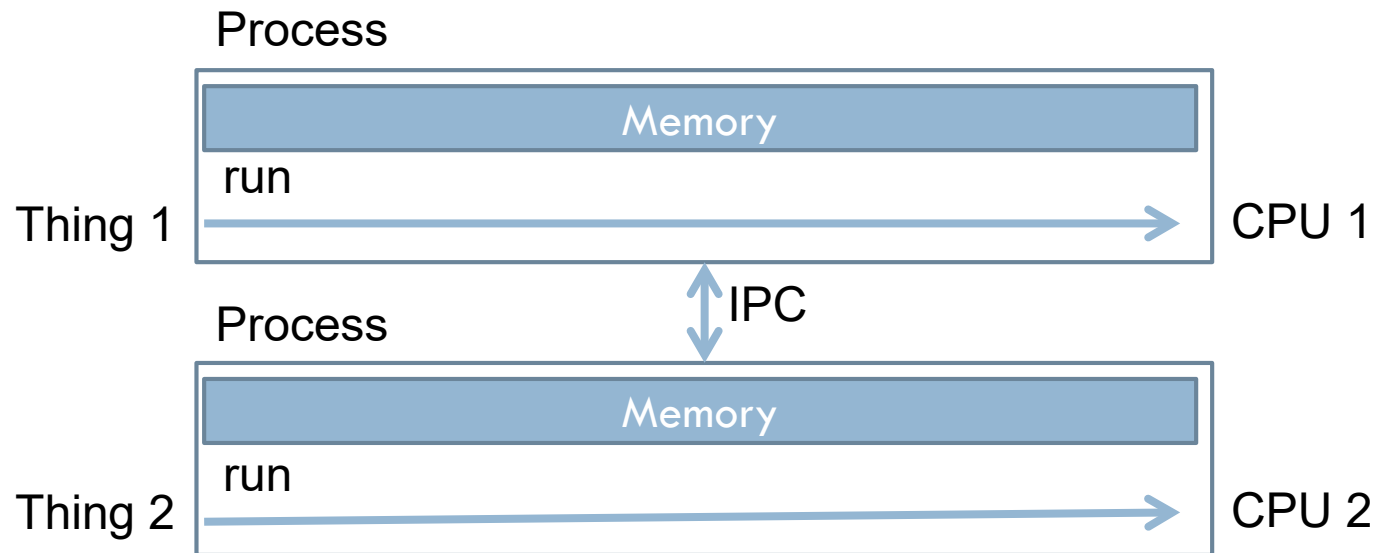
□ Simultaneous access to memory space

same process share memory, different processes share data



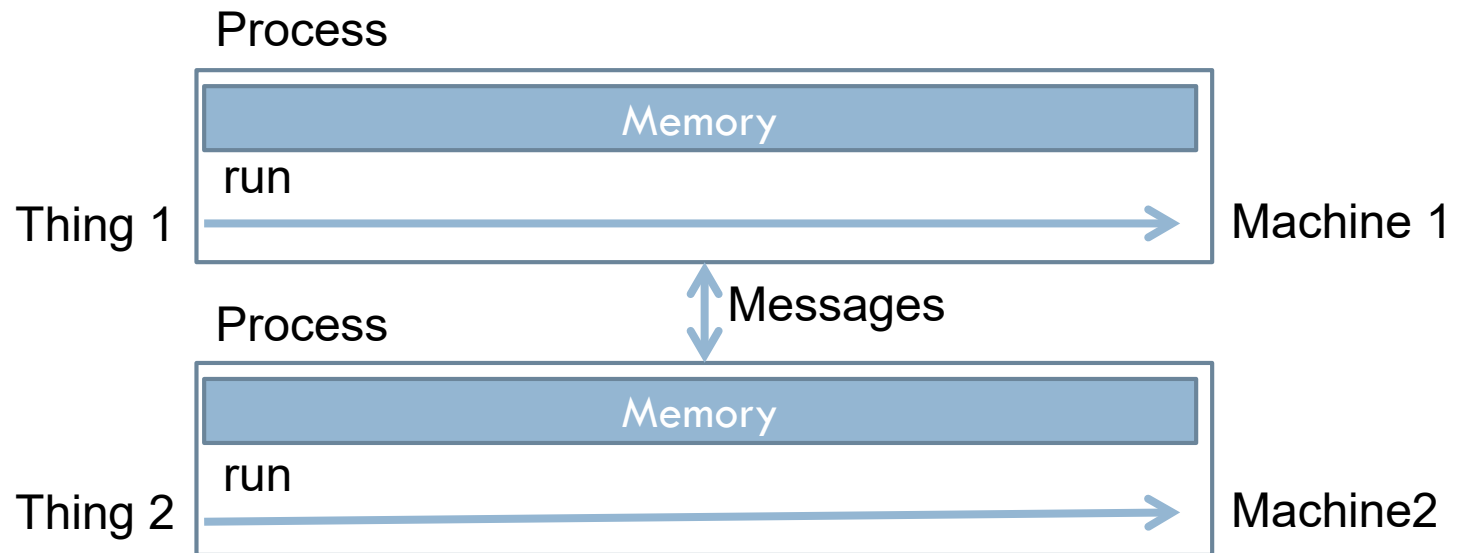
Separate processes, shared machine

- No shared memory
 - ▣ Inter-Process Communication (IPC)



Different machines, distributed computing

- Network
 - ▣ Sockets



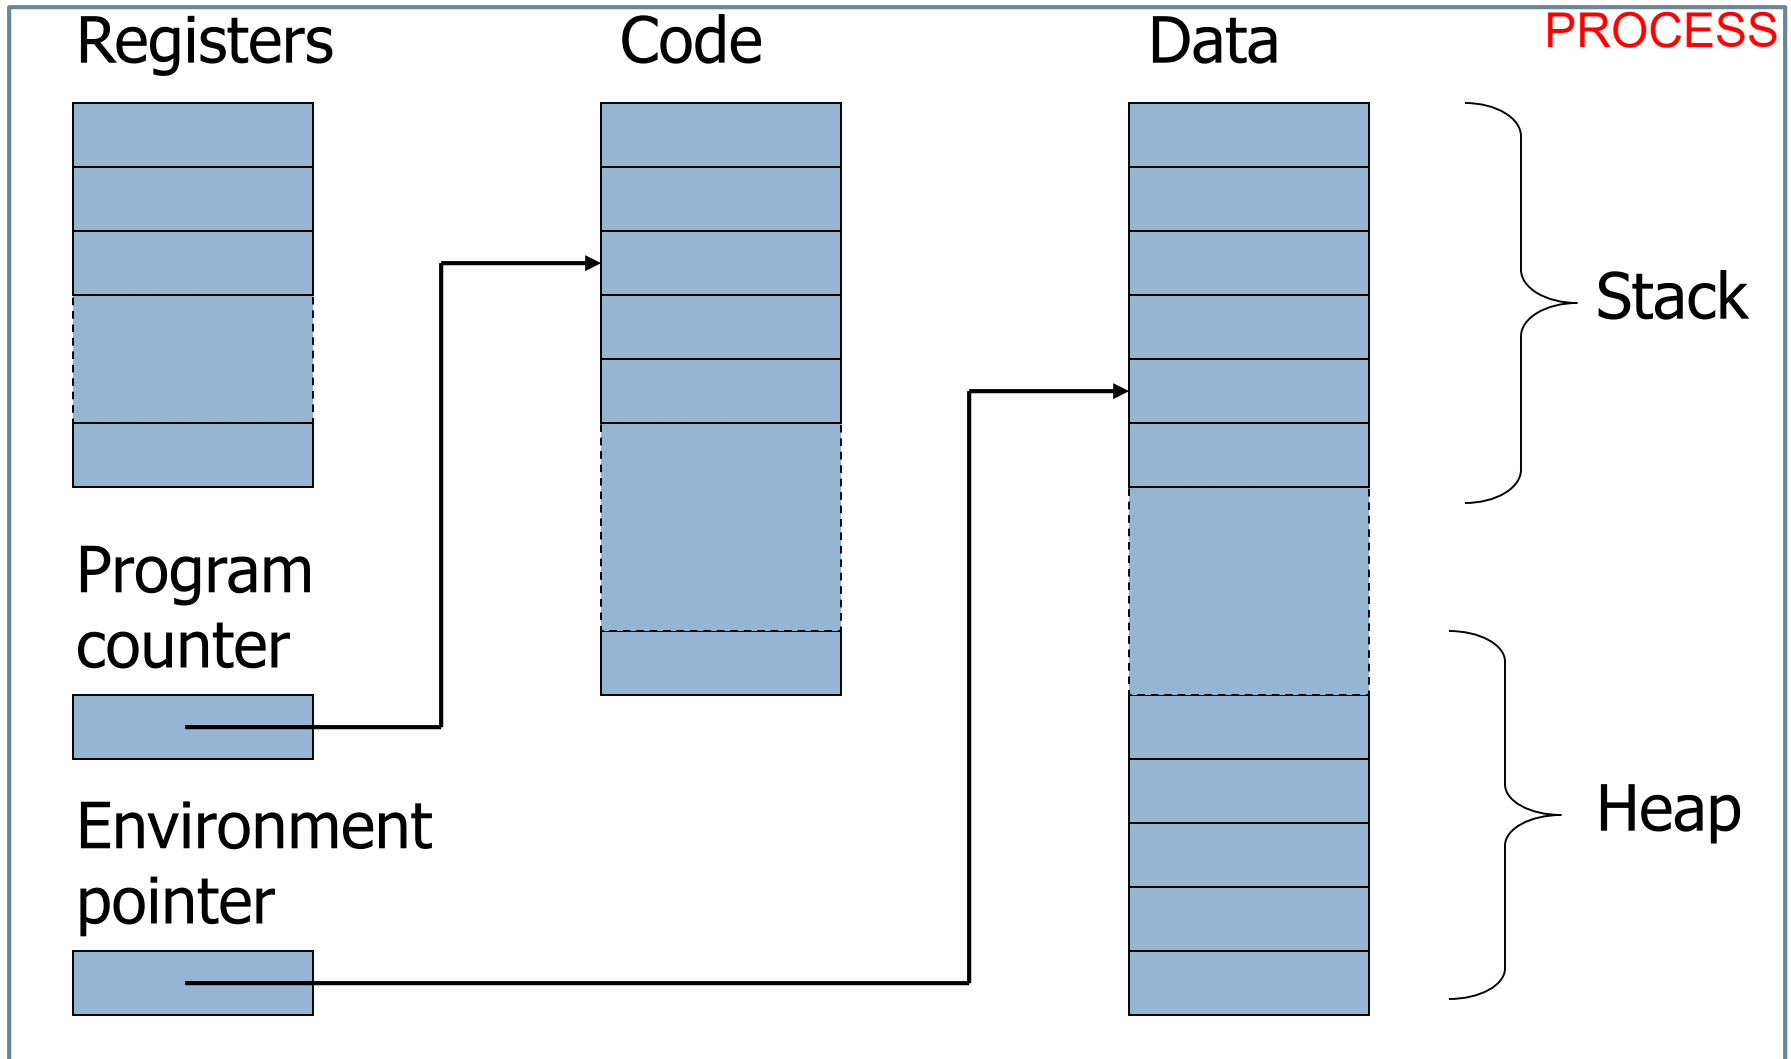


Shared Memory


Threads

Simplified Machine Model


slide 13



Threads



single-threaded



multithreaded

each thread has its
stack
heap is shared
files are shared

Threads: benefits

- Responsiveness
- Resource sharing
- [Performance in multi-core machines]

Threads: pitfalls

- ❑ Race conditions
- ❑ Heisenbugs happen once, can't control its appearance
- ❑ Easy to end up with a colossal code mess
- ❑ Heavy startup time for new threads create it is slow
 - ▣ [Thread pools to the rescue]

Two Types of Threads

17

- User-level threads libraries:

- POSIX *Pthreads*
- Java Threads
- .NET Threads
- Python Threads
- ...



This course

- Kernel threads




Thread API Examples

OOP Threads (Java/C#/PHP/Python)

v.1—instantiate Thread

```
class ThreadedObject implements Runnable {
    Thread t;
    ThreadedObject() {
        // Create a thread
        t = new Thread(this, "Demo Thread");
        System.out.println("Child thread: " + t);
        t.start(); // Start the thread
    }
    // This is the entry point for the ThreadedObject thread.
    public void run() {
        for (int i = 5; i > 0; i--) {
            System.out.println("Child Thread: " + i);
            // Let the thread sleep for a while.
            Thread.sleep(50);
        }
        System.out.println("Exiting child thread.");
    }
}
```

Thread runtime calls run



OOP Threads (Java/C#/PHP/Python)

v.1—instantiate Thread

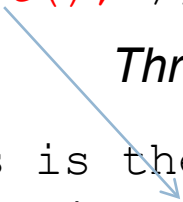
```
public class ThreadDemo {  
    public static void main(String args[]) {  
        new ThreadedObject(); // create a new thread  
        for (int i = 5; i > 0; i--) {  
            System.out.println("Main Thread: " + i);  
            Thread.sleep(100);  
        }  
        System.out.println("Main thread exiting.");  
    }  
}
```

OOP Threads (Java/C#/PHP/Python)

(v2—extend Thread)

```
class ThreadedObject extends Thread {
    ThreadedObject() {
        super("Demo Thread");
        System.out.println("Child thread: " + this);
        start(); // Start the thread
    }
    Thread runtime calls run

    // This is the entry point for the thread.
    public void run() {
        for (int i = 5; i > 0; i--) {
            System.out.println("Child Thread: " + i);
            // Let the thread sleep for a while.
            Thread.sleep(50);
        }
        System.out.println("Exiting child thread.");
    }
}
```



OOP Threads (Java/C#/PHP/Python)

(v2—extend Thread)

```
public class ThreadDemo {  
    public static void main(String args[]) {  
        new ThreadedObject(); // create a new thread  
        for (int i = 5; i > 0; i--) {  
            System.out.println("Main Thread: " + i);  
            Thread.sleep(100);  
        }  
        System.out.println("Main thread exiting.");  
    }  
}
```

Functional Threads

(C/C++/Perl/**Racket**)

```
#lang racket
```

```
(thread (lambda ()  
  (for ([i 10])  
    (sleep 2)  
    (printf "thread 1\n")))))
```

```
(thread (lambda ()  
  (for ([i 20])  
    (sleep 1)  
    (printf "thread 2\n")))))
```

Function to run



Functional Threads

(C/C++/Perl/Racket)

```
#include <string>
#include <iostream>
#include <thread>
using namespace std;
```

```
//The function we want to make the thread run.
```

```
void task1(string msg)
{
    cout << "task1 says: " << msg;
}
```

```
int main()
{
    // Constructs the new thread and runs it. Does not block
    thread t1(task1, "Hello");

    //Makes the main thread wait for the new thread to finish
    t1.join();
}
```

Function to run



Common operations on threads

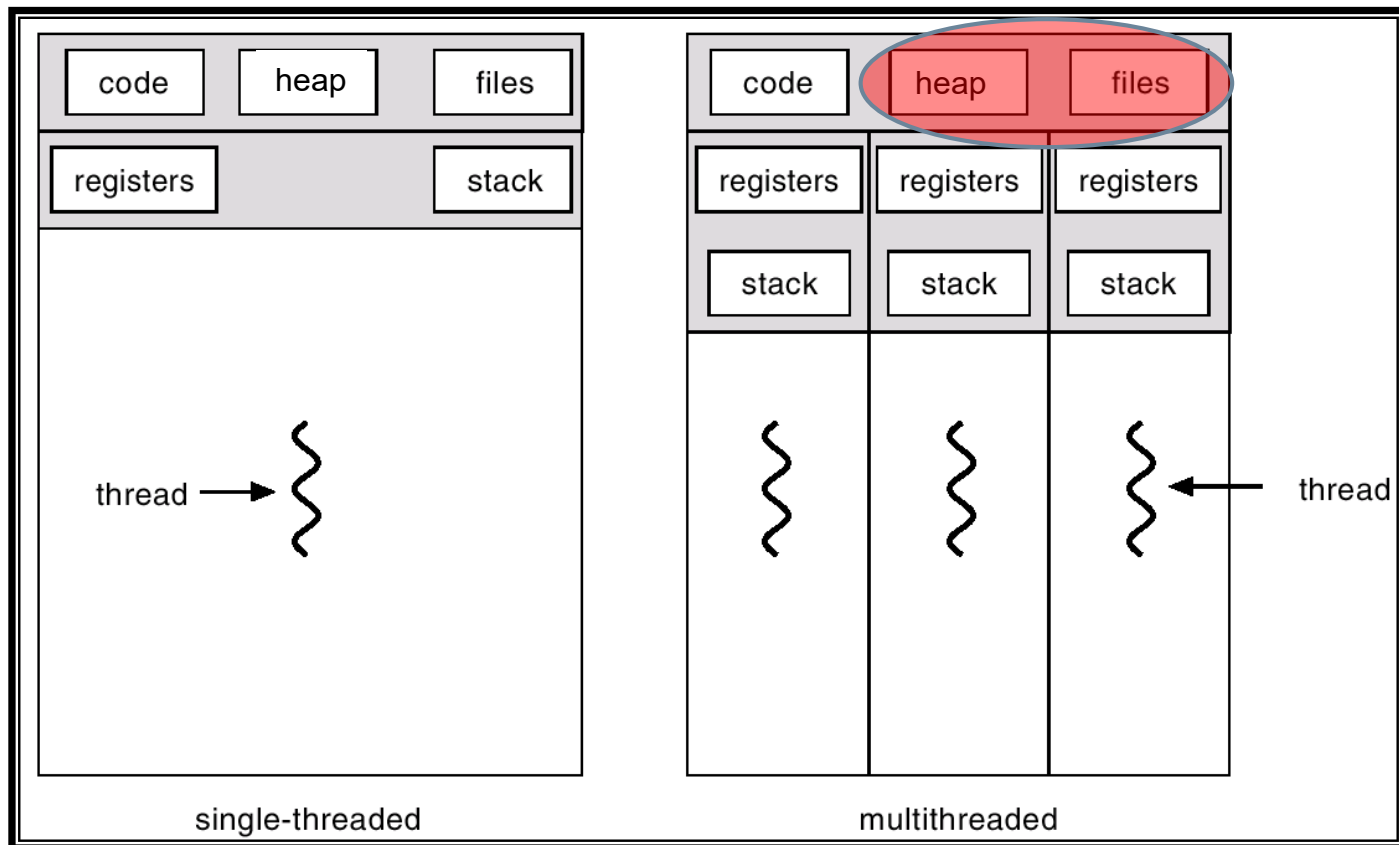
- Create
 - Start
 - Sleep
 - Suspend
 - Resume
 - Yield
 - Stop
-
- (you'll find most of these in most thread APIs)



Concurrency Control

Threads – shared resources

DANGER



Shared data problems

```
class ThreadedObject extends Thread {
    Counter c;
    ThreadedObject(Counter _c) {
        super("Demo Thread");
        System.out.println("Child thread: " + this);
        c = _c;
        start(); // Start the thread
    }

    public void run() {
        for (int i = 500; i > 0; i--) {
            System.out.println("Child Thread: " + i);
            if (i % 2 == 0) c.increment();
            else c.decrement();
            Thread.sleep(50);
        }
        System.out.println("Exiting child thread.");
    }
}
```

Shared data problems

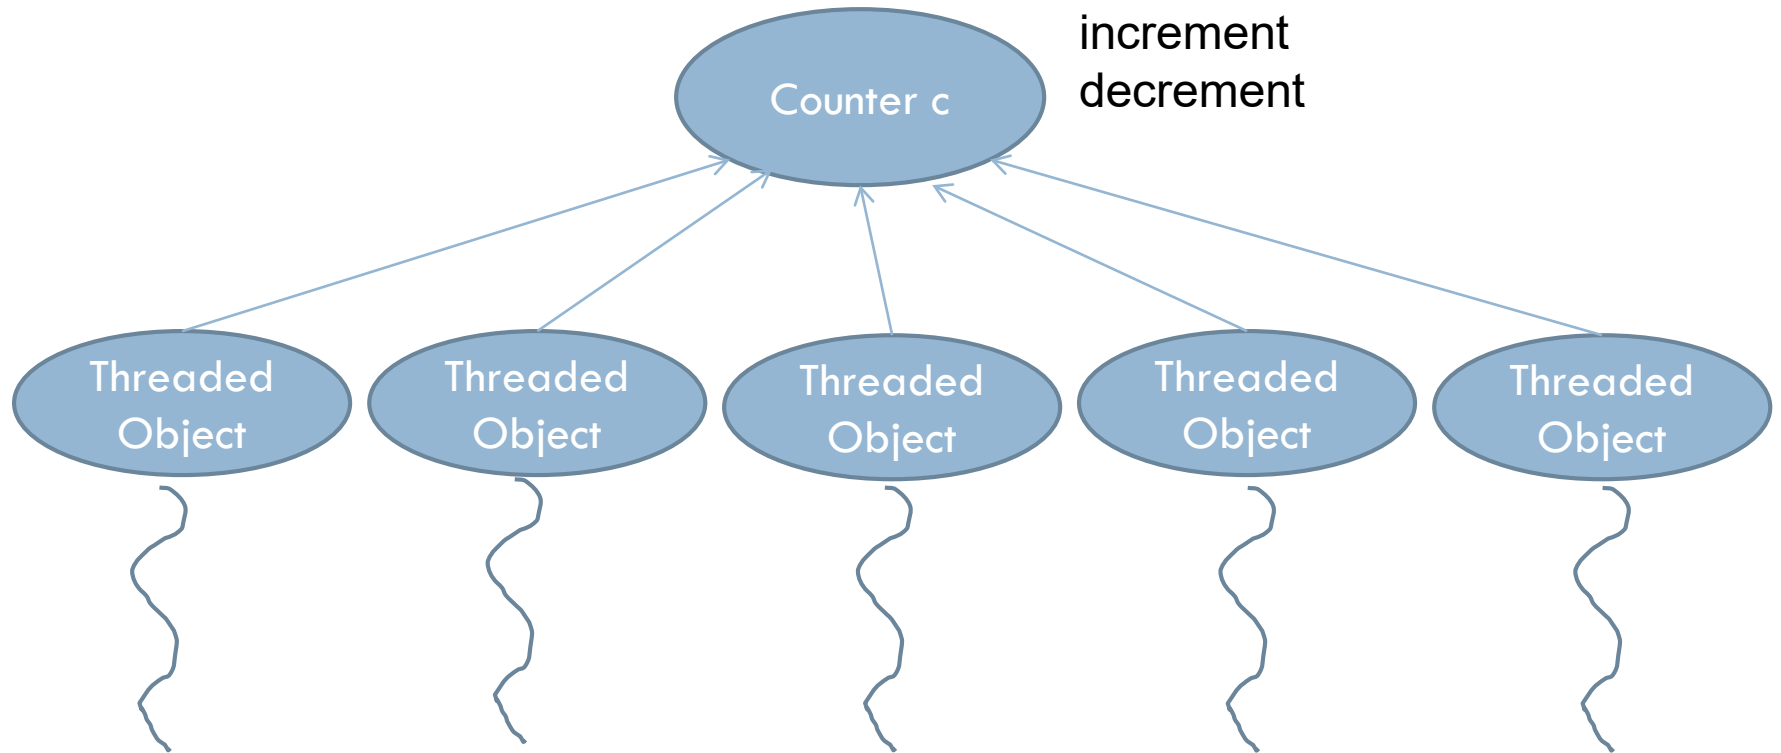
```
class Counter {  
    private int c = 0;  
  
    public void increment() {  
        c = c + 1;  
    }  
  
    public void decrement() {  
        c = c - 1;  
    }  
}
```

Shared data problems

```
public class ThreadDemo {
    public static void main(String args[]) {
        Counter c = new Counter();
        ThreadedObject[] ao = new ThreadedObject[5];
        // Create 5 threads
        for (int i = 5; i > 0; i--)
            ao[i] = new ThreadedObject(c); // create a new thread


        // Wait for the threads to finish
        for (ThreadedObject a : ao)
            a.join();
        System.out.println("Main thread exiting.");
    }
}
```

Shared data problems



Things that can go wrong

```
class Counter {  
    private int c = 0;  
  
    public void increment() {  
        c = c + 1;  
    }  
  
    public void decrement() {  
        c = c - 1;  
    }  
}
```



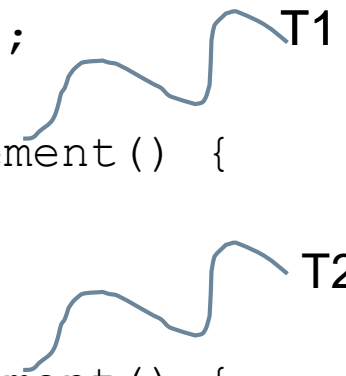
One action lost:

*T1 enters the method
T1 reads value of c (e.g. 3)
T1 is interrupted by scheduler
T2 enters the method
T2 reads value of c (3)
T2 adds 1 to c (4)
T2 assigns c to the new value (4)
T2 returns
T1 is resumed
T1 adds 1 to [old] value of c (4)
T1 assigns c to that value (4)*

Race condition

Things that can go wrong

```
class Counter {  
    private int c = 0;  
  
    public void increment() {  
        c = c + 1;  
    }  
  
    public void decrement() {  
        c = c - 1;  
    }  
}
```



One action lost:

T1 enters increment

T2 enters decrement

T1 reads value of c (e.g. 3)

T2 reads value of c (3)

T1 adds 1 to its value of c (4)

T2 subtracts 1 to its value of c (2)

T1 assigns new value to c (4)

T2 assigns c (2)

Race condition

Race conditions

- ❑ Corruption of shared data due to thread scheduling
- ❑ Very hard to deal with
 - ▣ Non-deterministic
 - ▣ Hard to reproduce
 - ▣ Program may be ok for long time until it hits a race condition

Concurrency control

- Primitives for controlling the execution of concurrent threads over the same code

Concurrency control (Java)

```
class Counter {  
    private int c = 0;  
  
    public synchronized void increment() {  
        c = c + 1;  
    }  
  
    public synchronized void decrement() {  
        c = c - 1;  
    }  
}
```

Only one thread at a time can enter the synchronized methods

Concurrency control (Java)

```
class Counter {  
    private int c = 0;  
    Object o = new Object();  
  
    public void increment() {  
        lock (this) {  
            c = c + 1;  
        }  
    }  
  
    public void decrement() {  
        lock (this) {  
            c = c - 1;  
        }  
    }  
}
```

Only one thread at a time can enter the locked blocks

Concurrency control (others)

```
class Counter {  
    private int c = 0;  
    Object o = new Object();  
  
    public void increment() {  
        lock (o) {  
            c = c + 1;  
        }  
    }  
  
    public void decrement() {  
        lock (o) {  
            c = c - 1;  
        }  
    }  
}
```

Only one thread at a time can enter the locked blocks

Concurrency control (others)

```
class Counter {  
    private int c = 0;  
    Lock o = new Lock();  
  
    public void increment() {  
        o.acquire();  
        c = c + 1;  
        o.release();  
    }  
  
    public void decrement() {  
        o.acquire();  
        c = c - 1;  
        o.release();  
    }  
}
```

Only one thread at a time can enter the locked blocks

Locks

- They seem easy
- They are **very hard** to manage

Lock management

- Acquired locks must always be released
 - ▣ Including when exceptions occur
 - ▣ A lock that is not released will prevent any other thread from entering the block forever

Over-locking → Deadlocks

```
class Counter {  
    private int c = 0;  
    Object o1 = new Object();  
    Object o2 = new Object();  
    public void increment() {  
        lock (o1) {  
            lock (o2) {  
                c = c + 1;  
            }  
        }  
    }  
  
    public void decrement() {  
        lock (o2) {  
            lock (o1) {  
                c = c - 1;  
            }  
        }  
    }  
}
```

The diagram illustrates a potential deadlock scenario. Two blue arrows originate from the text 'o1, then o2' and point to the 'lock (o1)' and 'lock (o2)' calls within the 'increment()' method. Another two blue arrows originate from the text 'o2, then o1' and point to the 'lock (o2)' and 'lock (o1)' calls within the 'decrement()' method. This shows that 'increment()' acquires o1 first, then o2, while 'decrement()' acquires o2 first, then o1, which can lead to a deadlock if both methods are executed simultaneously.

AVOID!

Deadlocks vs. Data Races

Too much concurrency control

deadlocks

Not enough concurrency control

data races



Threads Summary

- Nice, but...
 - ▣ Synchronization primitives are hard to get right
 - Deadlocks
 - Race conditions
 - ▣ Many corner cases
 - ▣ Bugs hard to reproduce
 - ▣ Potentially bad performance (locking is heavy)

- Many reasons to avoid threads!
 - ▣ Must use them under strict constraints